



kWh Energy Forecasting Project

1 - Importing Libraries

```
In [378...]:  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import sklearn  
import matplotlib.dates as mdates  
from datetime import datetime  
import prophet as Prophet  
from sklearn.model_selection import train_test_split  
from statsmodels.tsa.api import ARIMA, SARIMAX  
import pmdarima  
from statsmodels.tsa.stattools import adfuller  
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf  
from sklearn.tree import DecisionTreeRegressor, export_text  
from sklearn.metrics import mean_squared_error, root_mean_squared_error  
from sklearn.ensemble import RandomForestRegressor  
from statsmodels.tsa.holtwinters import ExponentialSmoothing
```

2 - Features

```
In [379...]:  
def features(df):  
    daily_data = pd.read_csv('daily - Sheet1.csv')  
  
    daily_data = daily_data.rename(columns={'Start': 'date'})  
  
    daily_data['date'] = (pd.to_datetime(daily_data['date'], format='%m/%d/%Y'))  
  
    daily_data['day'] = daily_data['date'].dt.day_name()  
    day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']  
    daily_data['day'] = pd.Categorical(daily_data['day'], categories=day_order)  
  
    daily_data['day_number'] = daily_data['date'].dt.day_of_week  
    weekdays = daily_data[daily_data['day_number'] < 5]  
    weekends = daily_data[daily_data['day_number'] >= 5]  
  
    daily_data['month'] = daily_data['date'].dt.month_name()  
    daily_data['month_number'] = daily_data['date'].dt.month # <-- Added this  
    daily_data['weekday'] = (daily_data['day_number'] < 5).astype(int)  
  
    daily_data = daily_data.set_index('date', inplace=False)  
    return daily_data
```

```
In [380...]: df = features('daily - Sheet1.csv')
```

```
df.head()
```

Out[380...]

	kWh	day	day_number	month	month_number	weekday
date						
2025-06-01	5.05	Sunday	6	June	6	0
2025-06-02	6.86	Monday	0	June	6	1
2025-06-03	9.10	Tuesday	1	June	6	1
2025-06-04	7.32	Wednesday	2	June	6	1
2025-06-05	12.88	Thursday	3	June	6	1

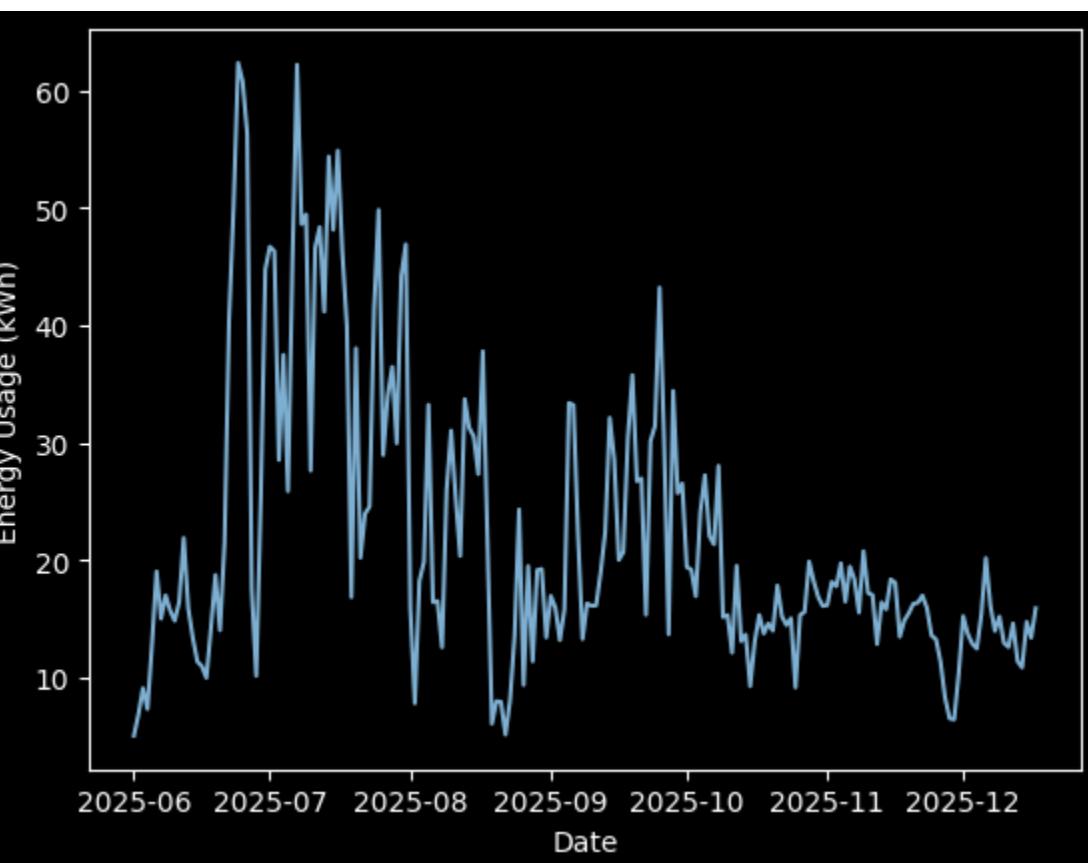
3 - Exploratory Data Analysis

3.1 - kWh Time Series Graph

In [386...]

```
sns.lineplot(df['kWh'], color = sns.color_palette()[4])
plt.xlabel('Date')
plt.ylabel('Energy Usage (kWh)')
plt.figure(figsize=(18,6))
```

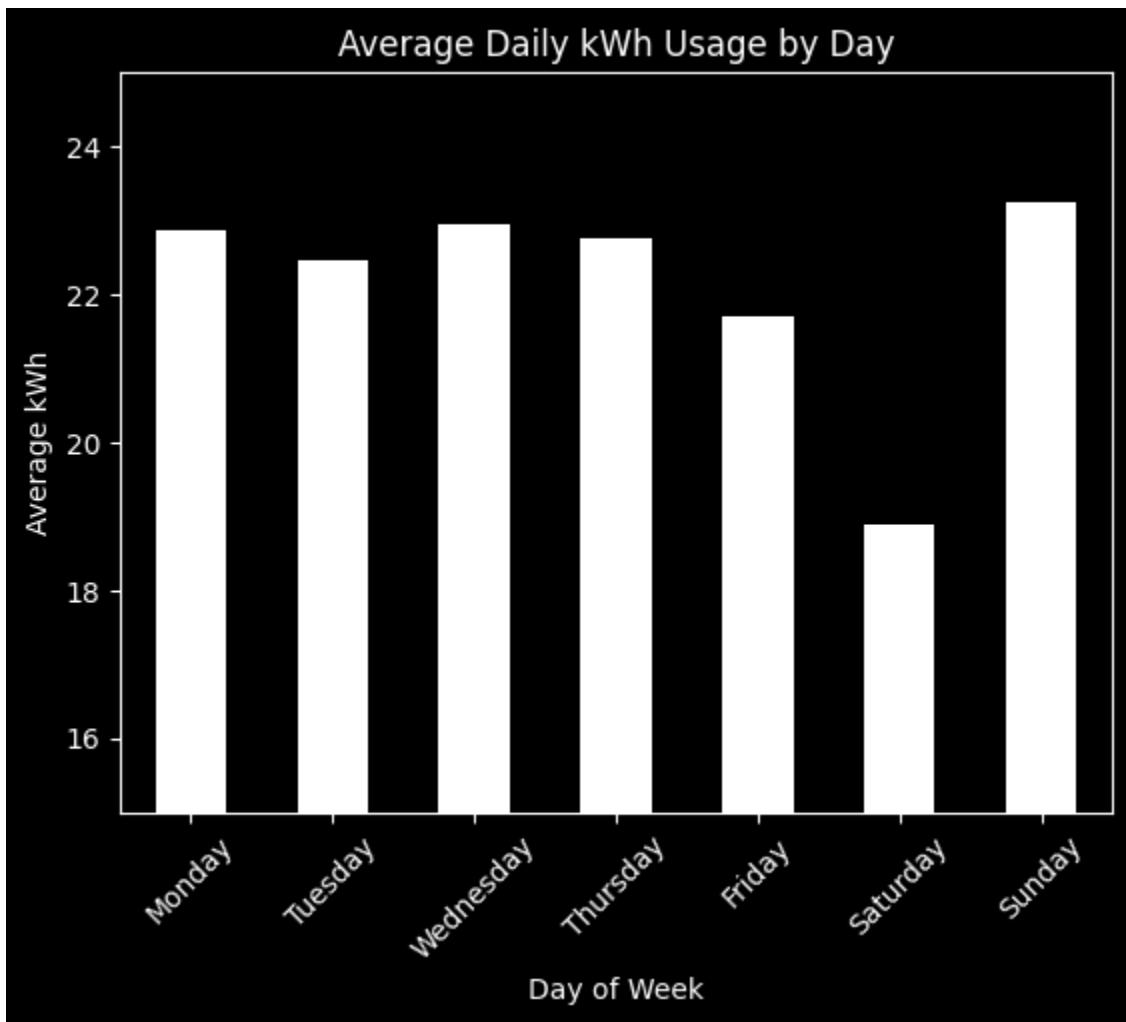
Out[386...]



<Figure size 1800x600 with 0 Axes>

3.2 - Average kWh Per Day

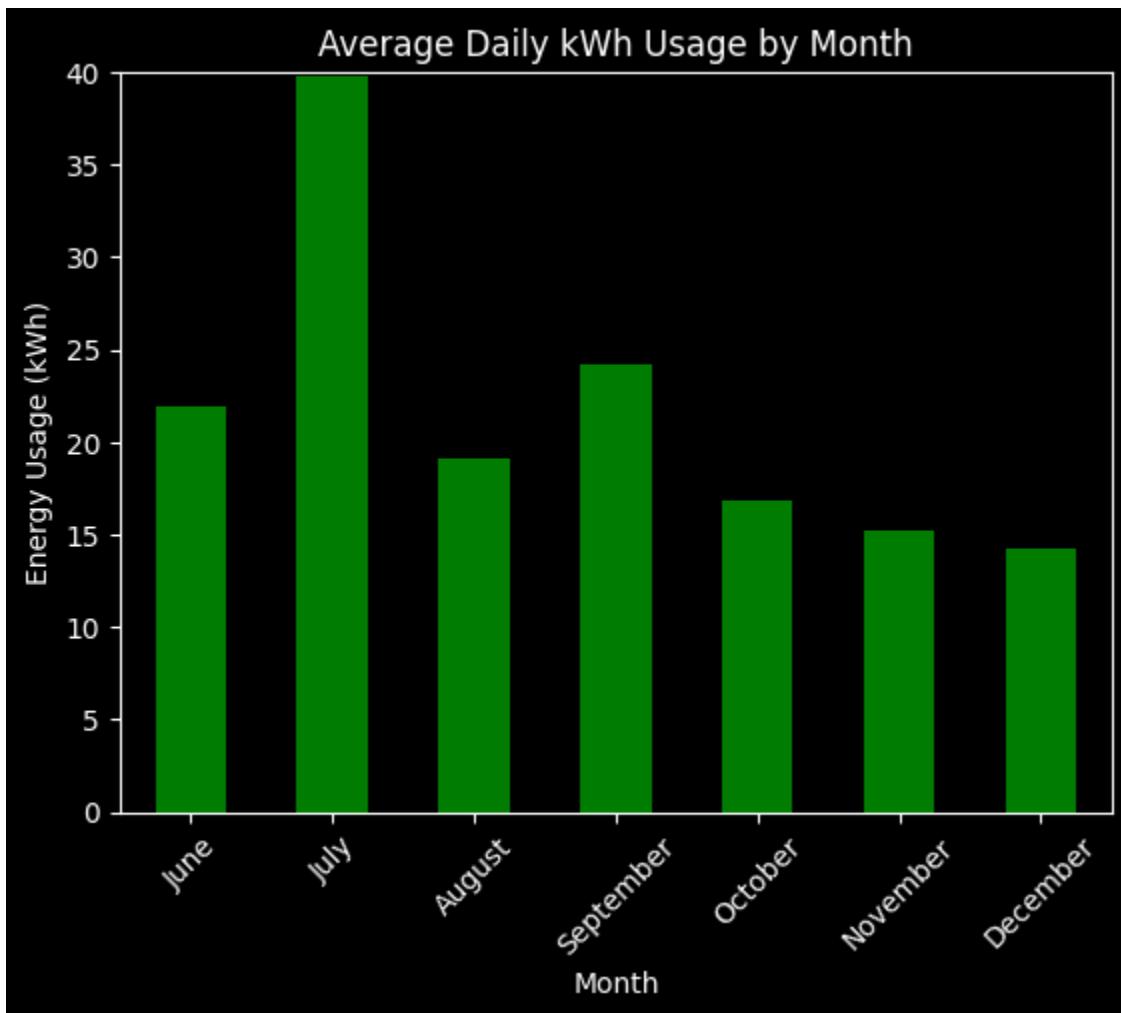
```
In [387...]: df.groupby('day', observed = True)[['kWh']].mean().plot(kind='bar', color='white')
plt.ylabel('Average kWh')
plt.xlabel('Day of Week')
plt.title('Average Daily kWh Usage by Day')
plt.xticks(rotation=45)
plt.ylim(15,25)
plt.show()
(df.groupby('day', observed = True)[['kWh']].mean())
```



```
Out[387...]: day
Monday      22.879310
Tuesday     22.468621
Wednesday   22.941379
Thursday    22.770357
Friday      21.705714
Saturday    18.884643
Sunday      23.262069
Name: kWh, dtype: float64
```

3.3 - Average kWh by Month

```
In [389...]: month_order = [  
    'June',  
    'July', 'August', 'September', 'October', 'November', 'December'  
]  
  
(  
    df.groupby('month')['kWh']  
        .mean()  
        .reindex(month_order)  
        .plot(kind='bar', color='Green')  
)  
  
plt.ylabel('Energy Usage (kWh)')  
plt.xlabel('Month')  
plt.title('Average Daily kWh Usage by Month')  
plt.xticks(rotation=45)  
plt.ylim(0, 40)  
plt.show()
```

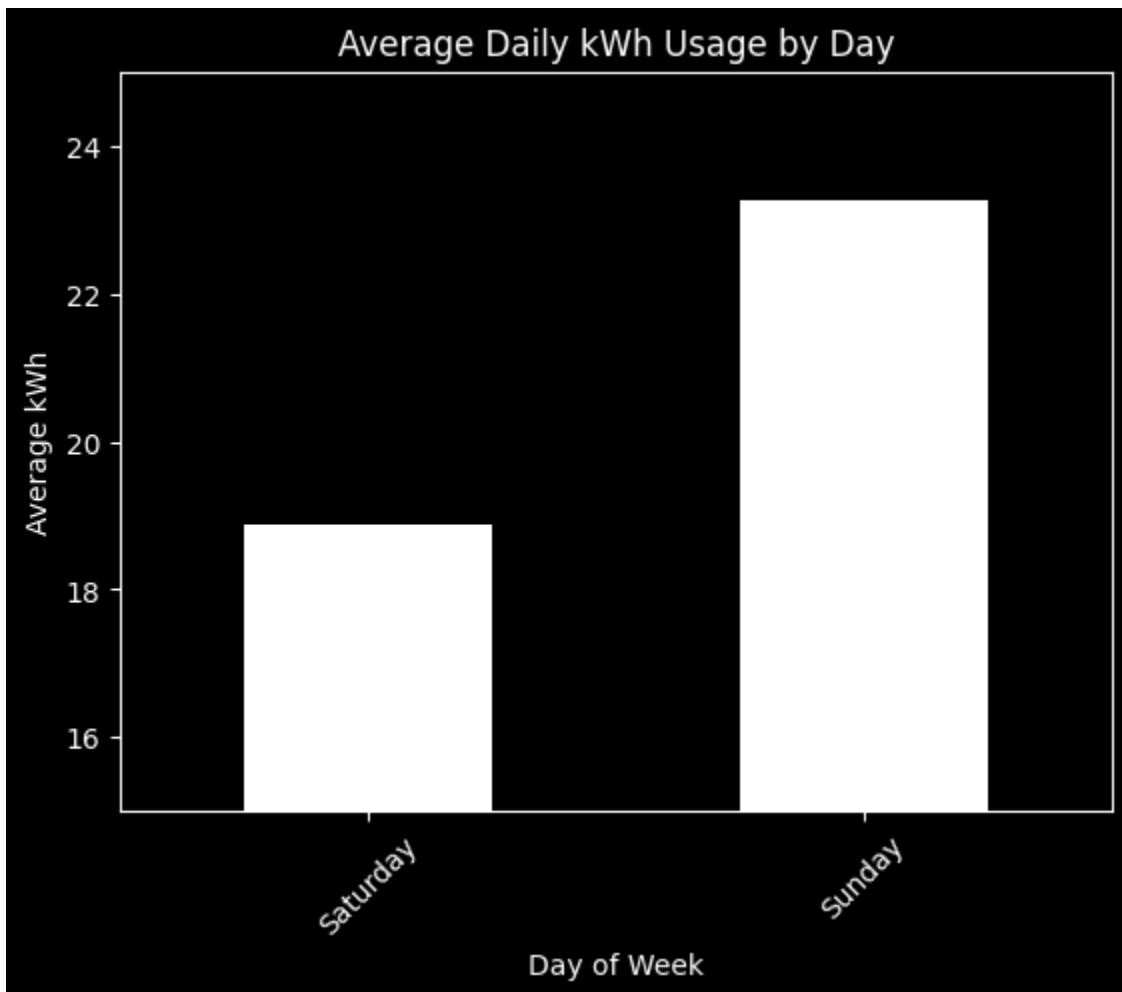


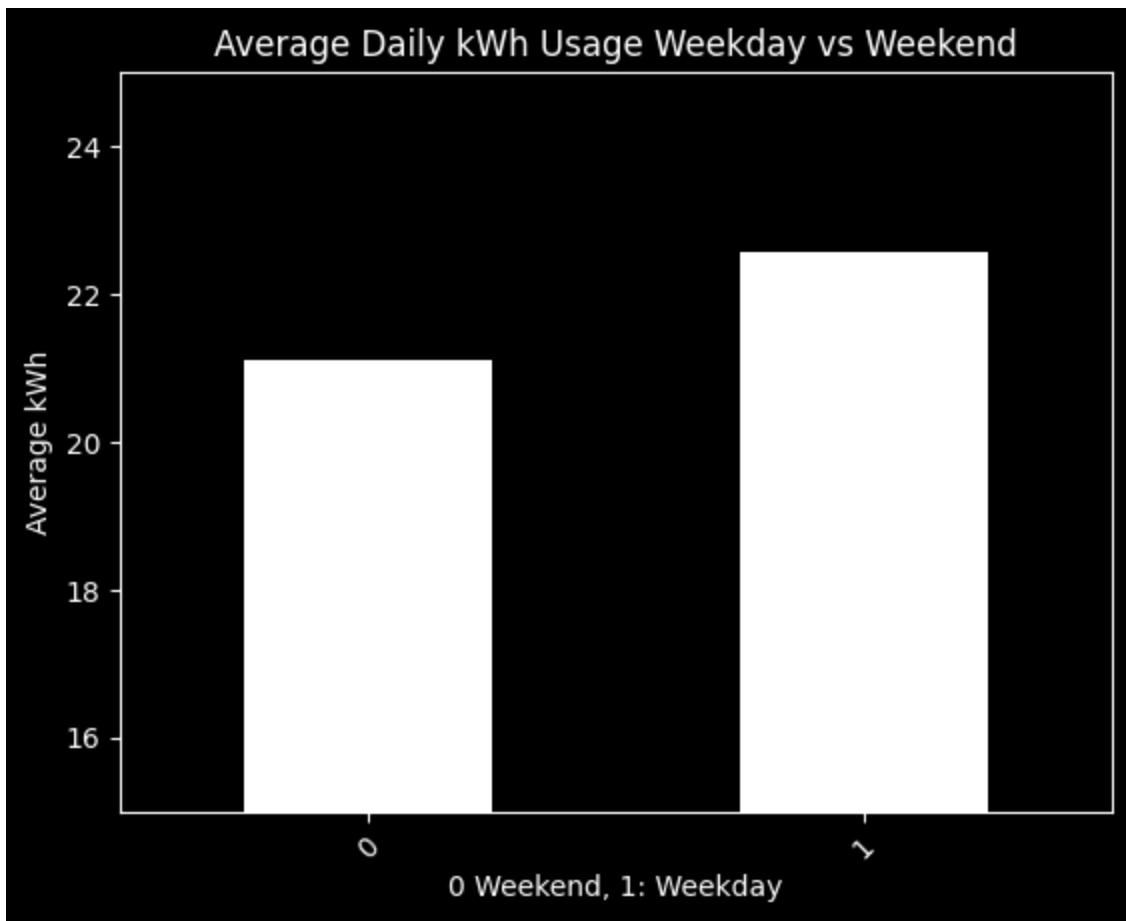
3.4 - Weekends vs Weekdays

In [418...]

```
weekends = df[df['day_number'] > 4]
weekends.groupby('day', observed = True)[['kWh']].mean().plot(kind='bar', color=
plt.ylabel('Average kWh')
plt.xlabel('Day of Week')
plt.title('Average Daily kWh Usage by Day')
plt.xticks(rotation=45)
plt.ylim(15,25)
plt.show()

(df.groupby('weekday', observed = True)[['kWh']].mean()).plot(kind='bar', color=
plt.ylabel('Average kWh')
plt.xlabel('0 Weekend, 1: Weekday')
plt.title('Average Daily kWh Usage Weekday vs Weekend')
plt.xticks(rotation=45)
plt.ylim(15,25)
plt.show()
```





4 - Train / Split

```
In [399...]: X = df.drop(columns=['kWh', 'day', 'month'])    # all real predictors
y = df['kWh'].values

split_point = int(len(X) * 0.8)

X_train = X.iloc[:split_point]
X_test = X.iloc[split_point:]

y_train = y[:split_point]
y_test = y[split_point:]
```

5 - Arima

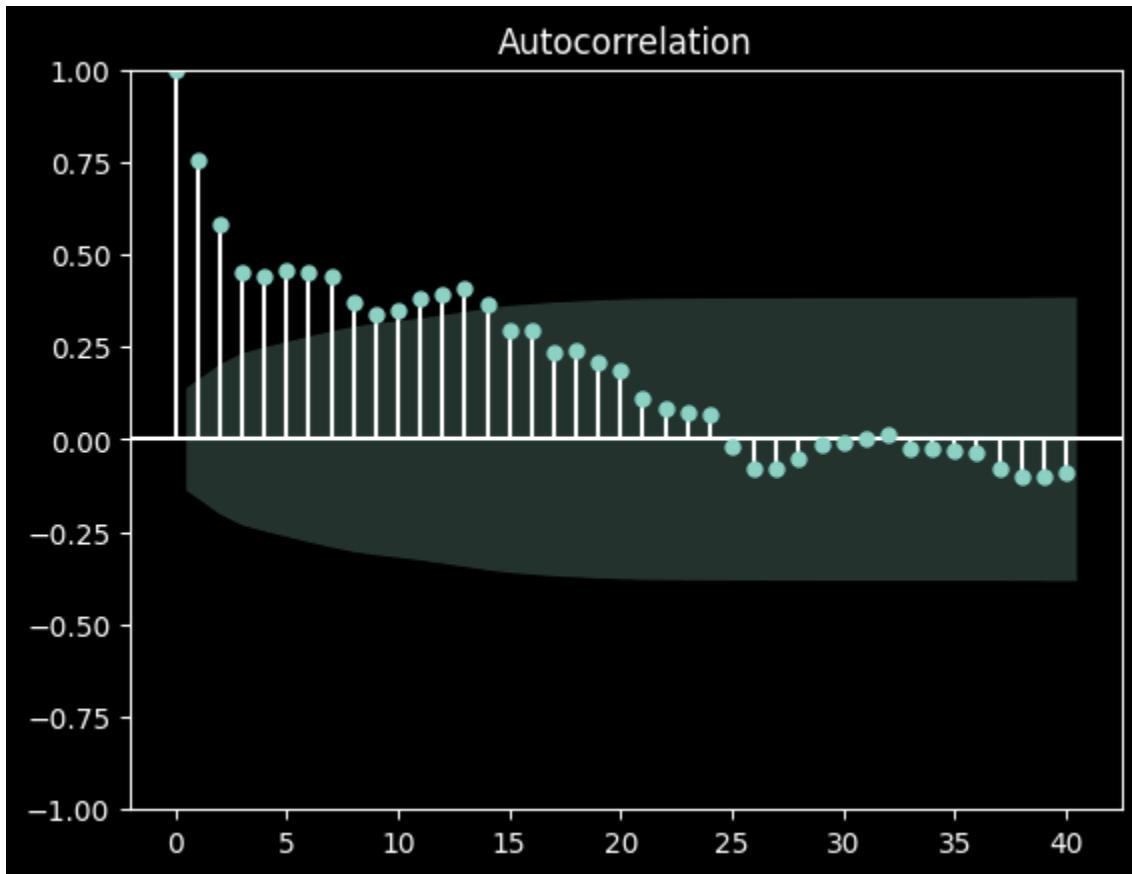
5.1 - ADF Test for Stationarity

```
In [400...]: adf_test = adfuller(df['kWh'].values)
# Output the results
print('ADF Statistic: %f' % adf_test[0])
print('p-value: %f' % adf_test[1])
```

```
ADF Statistic: -2.886002
p-value: 0.047000
```

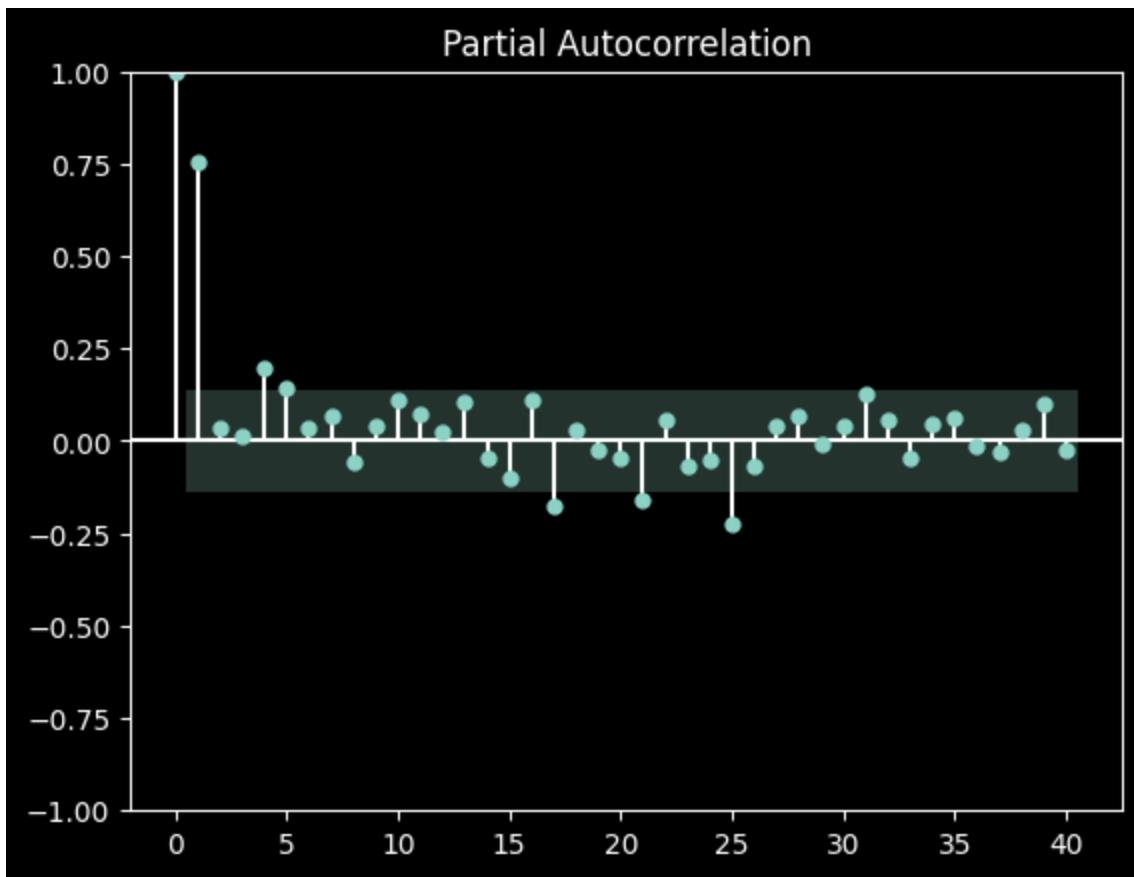
5.2 - Autocorrelation Function

```
In [401]: plot_acf(df['kWh'].values, lags=40)
plt.show()
```



5.3 - Partial Autocorrelation Function

```
In [402]: plot_pacf(df['kWh'].values, lags=40)
plt.show()
```



5.4 - Auto Arima

```
In [404]: from pmdarima import auto_arima

optimal_model = auto_arima(y_train, start_p=0, start_q=0,
                           test='adf',
                           max_p=3, max_q=3,
                           m=1,
                           d=None,
                           seasonal=True,
                           start_P=0,
                           D=0,
                           trace=True,
                           error_action='ignore',
                           suppress_warnings=True,
                           stepwise=True)

print(optimal_model.summary())
```

Performing stepwise search to minimize aic

ARIMA(0,1,0)(0,0,0)[0]	intercept	:	AIC=1173.051, Time=0.04 sec
ARIMA(1,1,0)(0,0,0)[0]	intercept	:	AIC=1171.196, Time=0.01 sec
ARIMA(0,1,1)(0,0,0)[0]	intercept	:	AIC=1169.277, Time=0.01 sec
ARIMA(0,1,0)(0,0,0)[0]		:	AIC=1171.063, Time=0.00 sec
ARIMA(1,1,1)(0,0,0)[0]	intercept	:	AIC=1153.876, Time=0.03 sec
ARIMA(2,1,1)(0,0,0)[0]	intercept	:	AIC=1154.521, Time=0.03 sec
ARIMA(1,1,2)(0,0,0)[0]	intercept	:	AIC=1155.066, Time=0.02 sec
ARIMA(0,1,2)(0,0,0)[0]	intercept	:	AIC=1158.878, Time=0.01 sec
ARIMA(2,1,0)(0,0,0)[0]	intercept	:	AIC=1171.193, Time=0.01 sec
ARIMA(2,1,2)(0,0,0)[0]	intercept	:	AIC=inf, Time=0.08 sec
ARIMA(1,1,1)(0,0,0)[0]		:	AIC=1151.910, Time=0.02 sec
ARIMA(0,1,1)(0,0,0)[0]		:	AIC=1167.299, Time=0.01 sec
ARIMA(1,1,0)(0,0,0)[0]		:	AIC=1169.212, Time=0.01 sec
ARIMA(2,1,1)(0,0,0)[0]		:	AIC=1152.564, Time=0.02 sec
ARIMA(1,1,2)(0,0,0)[0]		:	AIC=1153.105, Time=0.01 sec
ARIMA(0,1,2)(0,0,0)[0]		:	AIC=1156.930, Time=0.01 sec
ARIMA(2,1,0)(0,0,0)[0]		:	AIC=1169.212, Time=0.01 sec
ARIMA(2,1,2)(0,0,0)[0]		:	AIC=inf, Time=0.06 sec

Best model: ARIMA(1,1,1)(0,0,0)[0]

Total fit time: 0.399 seconds

SARIMAX Results

```
=====
Dep. Variable:                      y      No. Observations:                 160
Model:                SARIMAX(1, 1, 1)   Log Likelihood:            -572.955
Date:                Sun, 21 Dec 2025   AIC:                         1151.910
Time:                       15:40:50     BIC:                         1161.117
Sample:                           0 - 160   HQIC:                        1155.649
Covariance Type:                  opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.5345	0.081	6.603	0.000	0.376	0.693
ma.L1	-0.8713	0.047	-18.727	0.000	-0.963	-0.780
sigma2	78.7210	7.327	10.744	0.000	64.361	93.082

=====

Ljung-Box (L1) (Q): 0.29 Jarque-Bera (JB): 1

2.29

Prob(Q): 0.59 Prob(JB):

0.00

Heteroskedasticity (H): 0.27 Skew:

-0.33

Prob(H) (two-sided): 0.00 Kurtosis:

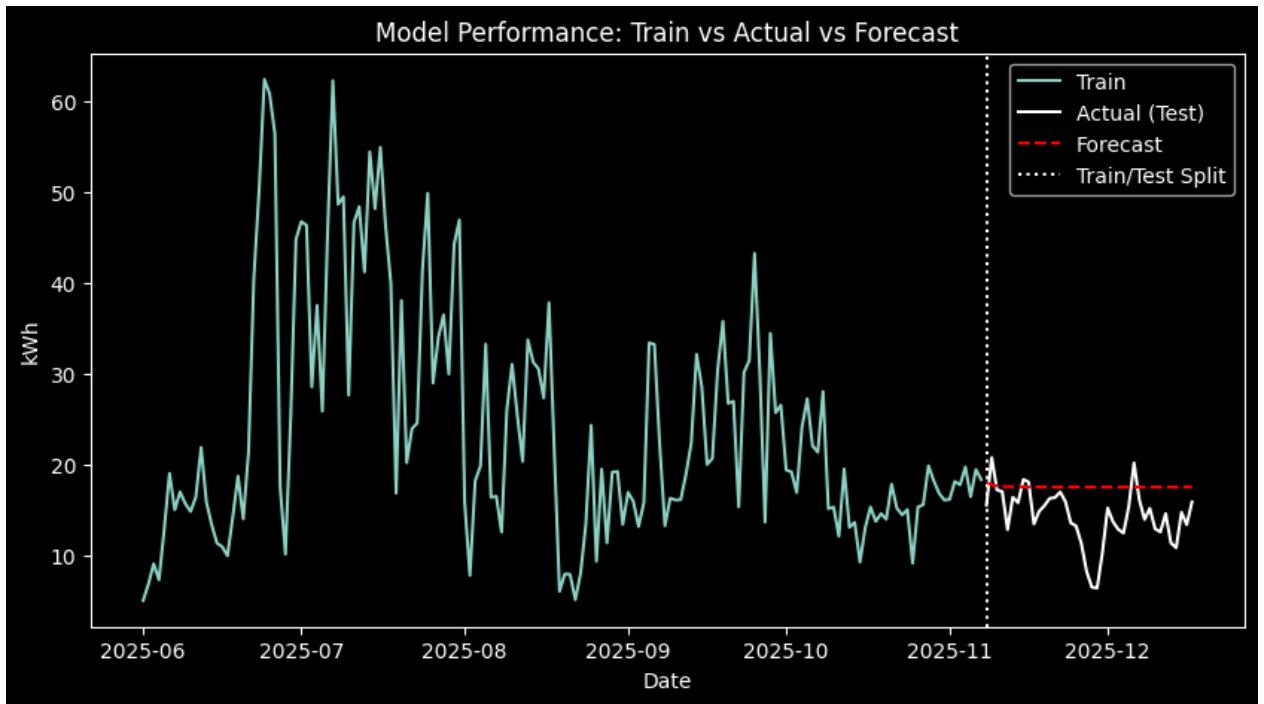
4.19

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [42]:  
import warnings  
warnings.filterwarnings("ignore", category=UserWarning, module="statsmodels.ts  
model = SARIMAX(y_train, order=(1,1, 1))  
model_fit = model.fit()
```

```
In [43]:  
n_test = len(y_test)  
forecast_values = model_fit.forecast(steps=n_test)  
  
test_index = df.index[-n_test:]  
  
forecast = pd.Series(forecast_values, index=test_index)  
y_test_series = pd.Series(y_test, index=test_index)  
  
plt.figure(figsize=(10, 5))  
  
plt.plot(df.index[:-n_test], df['kWh'].iloc[:-n_test], label='Train')  
plt.plot(test_index, y_test_series, label='Actual (Test)', color = 'white')  
plt.plot(test_index, forecast, label='Forecast', linestyle='--', color = 'red')  
  
plt.axvline(df.index[-n_test], linestyle=':', label='Train/Test Split')  
  
plt.title('Model Performance: Train vs Actual vs Forecast')  
plt.ylabel('kWh')  
plt.xlabel('Date')  
plt.legend()  
plt.show()  
  
  
from sklearn.metrics import mean_absolute_error, root_mean_squared_error, r2_s  
print("Test RMSE:", root_mean_squared_error(y_test_series, forecast))  
print("Test MAE:", mean_absolute_error(y_test_series, forecast))  
  
naive_forecast = np.repeat(y_train[-1], len(y_test))  
  
print("Naive RMSE:",  
      root_mean_squared_error(y_test, naive_forecast))
```



Test RMSE: 4.459848125927703

Test MAE: 3.592938531099827

Naive RMSE: 5.07309939780407

6.0 Regression Trees

```
In [429...]: y_train_series = pd.Series(y_train, index=df.index[:split_point])
y_test_series = pd.Series(y_test, index=df.index[split_point:])

regressor = DecisionTreeRegressor(max_depth=4, random_state=42)
regressor.fit(X_train, y_train_series)
y_pred = regressor.predict(X_test)

mse = root_mean_squared_error(y_test, y_pred)
print(f"Root Mean Squared Error: {mse:.4f}")

plt.figure(figsize=(12,6))

plt.plot(y_train_series.index, y_train_series, label='Train', color='blue')

plt.plot(y_test_series.index, y_test_series, label='Actual (Test)', color='White')

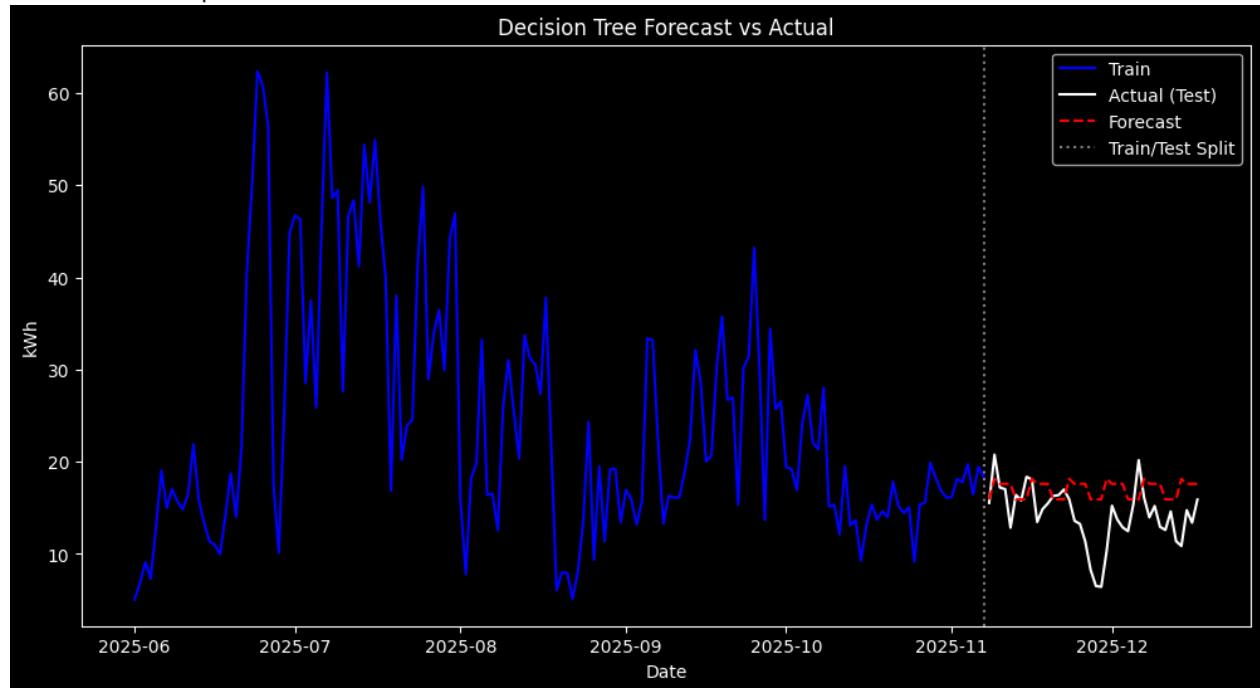
plt.plot(y_test_series.index, y_pred, label='Forecast', color='red', linestyle='dashed')

plt.axvline(y_train_series.index[-1], color='gray', linestyle=':', label='Train/Test Split')

plt.title('Decision Tree Forecast vs Actual')
```

```
plt.xlabel('Date')
plt.ylabel('kWh')
plt.legend()
plt.show()
```

Root Mean Squared Error: 4.1153



7.0 - Random Forrest

```
In [424]: rf_regressor = RandomForestRegressor(
    n_estimators=100,
    max_depth=4,
    random_state=42
)
rf_regressor.fit(X_train, y_train_series)

y_pred = rf_regressor.predict(X_test)

#Metrics
rmse = root_mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
print(f"Random Forest RMSE: {rmse:.4f}")
print(f"Random Forest MAE: {mae:.4f}")

feature_importance = pd.Series(
    rf_regressor.feature_importances_,
    index=X.columns
).sort_values(ascending=False)
print("\nFeature Importance:")
print(feature_importance)

plt.figure(figsize=(12,6))
```

```

plt.plot(y_train_series.index, y_train_series, label='Train', color='blue')
plt.plot(y_test_series.index, y_test_series, label='Actual (Test)', color='white')
plt.plot(y_test_series.index, y_pred, label='Forecast', color='red', linestyle='solid')
plt.axvline(y_train_series.index[-1], color='gray', linestyle='dotted', label='Train/Test Split')

plt.title('Random Forest Forecast vs Actual')
plt.xlabel('Date')
plt.ylabel('kWh')
plt.legend()
plt.show()

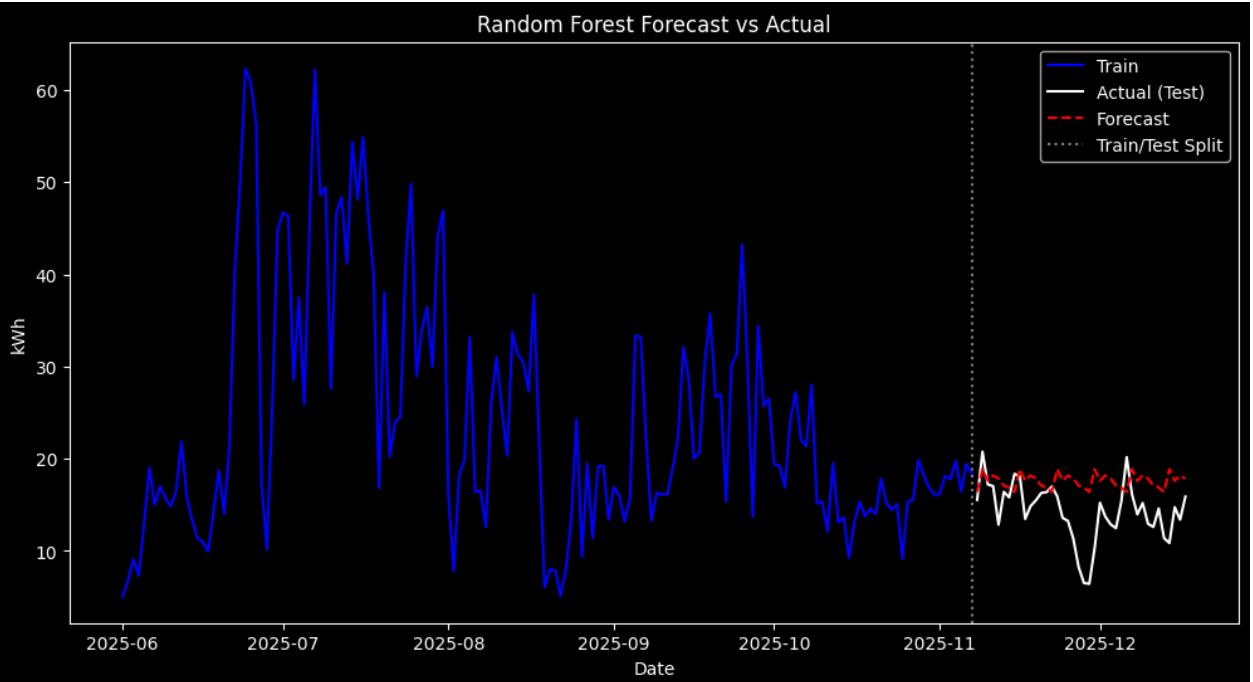
```

Random Forest RMSE: 4.5074

Random Forest MAE: 3.6762

Feature Importance:

month_number	0.771722
day_number	0.215202
weekday	0.013076
dtype:	float64



8.0 - Prophet

```

In [425]: prophet_df = df.reset_index()[['date', 'kWh']].rename(columns={'date': 'ds', 'kWh': 'y'})

split_point = int(len(prophet_df) * 0.8)
train_df = prophet_df.iloc[:split_point]
test_df = prophet_df.iloc[split_point:]

model = Prophet(daily_seasonality=False, weekly_seasonality=True, yearly_seasoc

```

```

model.fit(train_df)

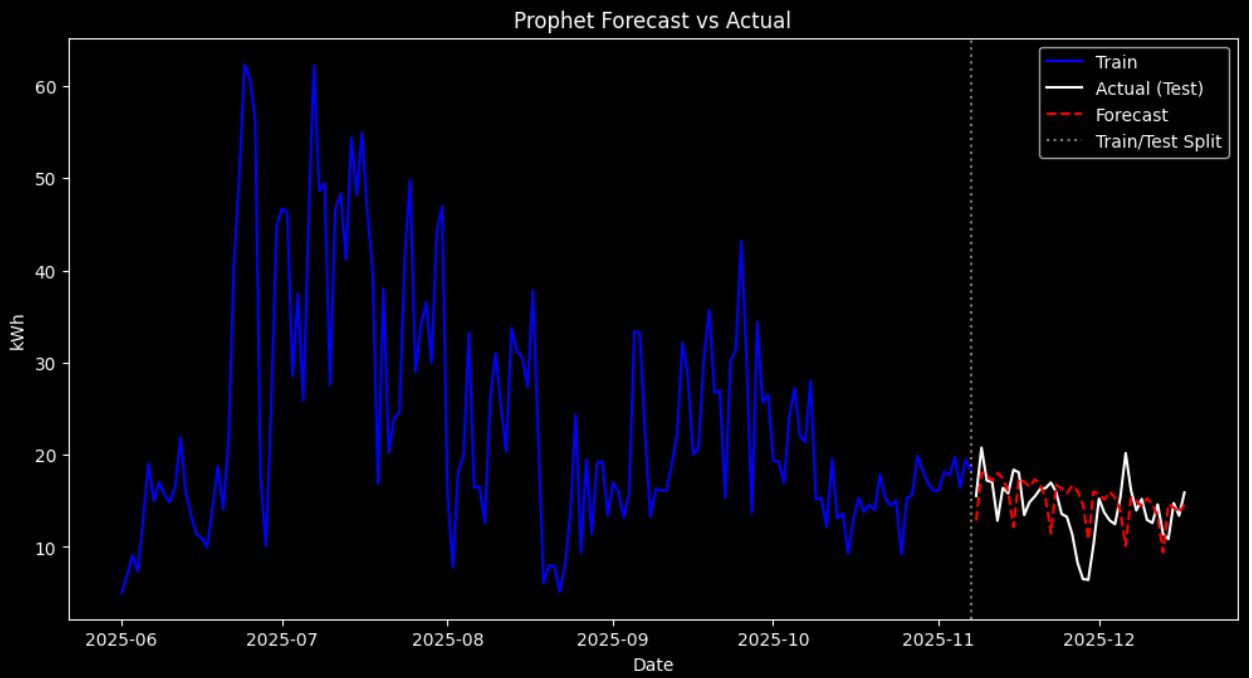
future = test_df[['ds']].copy()
forecast = model.predict(future)

plt.figure(figsize=(12,6))
plt.plot(train_df['ds'], train_df['y'], label='Train', color='blue')
plt.plot(test_df['ds'], test_df['y'], label='Actual (Test)', color='white')
plt.plot(forecast['ds'], forecast['yhat'], label='Forecast', color='red', line)
plt.axvline(train_df['ds'].iloc[-1], color='gray', linestyle=':', label='Train/Test Split')
plt.title('Prophet Forecast vs Actual')
plt.xlabel('Date')
plt.ylabel('kWh')
plt.legend()
plt.show()

#Metrics
rmse = root_mean_squared_error(test_df['y'], forecast['yhat'])
mae = mean_absolute_error(test_df['y'], forecast['yhat'])
print(f"Prophet RMSE: {rmse:.4f}")
print(f"Prophet MAE: {mae:.4f}")

```

15:50:01 - cmdstanpy - INFO - Chain [1] start processing
 15:50:01 - cmdstanpy - INFO - Chain [1] done processing



Prophet RMSE: 3.5739
 Prophet MAE: 2.6618

9.0 - Holt Winters Exponential Smoothing

In [428...]: `y_train_series = y_train_series.asfreq('D')`

```

y_test_series = y_test_series.asfreq('D')

from statsmodels.tools.sm_exceptions import ConvergenceWarning
warnings.filterwarnings("ignore", category=ConvergenceWarning)

hw_model = ExponentialSmoothing(
    y_train_series,
    trend='multiplicative',
    seasonal='additive',
    seasonal_periods=7
).fit()

y_pred = hw_model.forecast(len(y_test_series))

rmse = root_mean_squared_error(y_test_series, y_pred)
mae = mean_absolute_error(y_test_series, y_pred)
print(f"Holt-Winters RMSE: {rmse:.4f}")
print(f"Holt-Winters MAE: {mae:.4f}")

plt.figure(figsize=(12,6))
plt.plot(y_train_series.index, y_train_series, label='Train', color='blue')
plt.plot(y_test_series.index, y_test_series, label='Actual (Test)', color='white')
plt.plot(y_test_series.index, y_pred, label='Forecast', color='red', linestyle='dashed')
plt.axvline(y_train_series.index[-1], color='gray', linestyle=':', label='Train/Test Split')
plt.title('Holt-Winters Forecast vs Actual')
plt.xlabel('Date')
plt.ylabel('kWh')
plt.legend()
plt.show()

```

Holt-Winters RMSE: 3.5696

Holt-Winters MAE: 2.5214

